

Senior AI & Deep Learning Engineer

Email: adrian.vance@email.com Phone: +1 (555) 019-8811 Location: San Francisco, CA GitHub: github.com/avance-ai

PROFESSIONAL SUMMARY

Innovative AI Engineer with 6+ years of experience building, fine-tuning, and deploying deep learning models, LLMs, and computer vision systems into production scalable environments.

CORE COMPETENCIES

Frameworks: PyTorch, TensorFlow, Hugging Face, OpenCV, LangChain, Ray
MLOps & Deployment: FastAPI, Docker, Triton Inference Server, AWS SageMaker, MLflow, Kubernetes
Specializations: Large Language Models (LLMs), Computer Vision, Generative AI, Model Optimization (ONNX, TensorRT)

PROFESSIONAL EXPERIENCE

Lead AI Engineer | NeuralMind Systems

2022 – Present

- Fine-tuned open-source LLMs using LoRA/QLoRA for enterprise knowledge retrieval, improving answer precision by 28%.
- Deployed Triton Inference Server clusters handling 1M+ daily real-time inference requests with sub-50ms latency.
- Built automated MLOps retraining pipelines with MLflow and Airflow, cutting deployment prep time by 60%.

Computer Vision Engineer | VisionTech AI

2019 – 2022

- Developed YOLOv8 object detection pipelines for automated quality inspection, achieving 98.4% detection accuracy.
- Accelerated PyTorch model inference speed by 3.2x through TensorRT optimization.

EDUCATION

Ph.D. in Computer Science (Artificial Intelligence) - Stanford University

2015 – 2019

Certifications: AWS Certified Machine Learning – Specialty